

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv('HWs/HW3/Files/homework-3/data/titanic.csv')
```

```
In [ ]: data.head()
```

```
Out[ ]:
```

| | passenger_id | survived | pclass | name | sex | age | sib_sp | parch | ticket | fare | cabin | emba |
|---|--------------|----------|--------|---|--------|------|--------|-------|------------------|---------|-------|------|
| 0 | 1 | No | 3 | Braund, Mr. Owen Harris | male | 22.0 | 1 | 0 | A/5 21171 | 7.2500 | NaN | |
| 1 | 2 | Yes | 1 | Cumings, Mrs. John Bradley (Florence Briggs Th... | female | 38.0 | 1 | 0 | PC 17599 | 71.2833 | C85 | |
| 2 | 3 | Yes | 3 | Heikkinen, Miss. Laina | female | 26.0 | 0 | 0 | STON/O2. 3101282 | 7.9250 | NaN | |
| 3 | 4 | Yes | 1 | Futrelle, Mrs. Jacques Heath (Lily May Peel) | female | 35.0 | 1 | 0 | 113803 | 53.1000 | C123 | |
| 4 | 5 | No | 3 | Allen, Mr. William Henry | male | 35.0 | 0 | 0 | 373450 | 8.0500 | NaN | |

```
In [ ]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   passenger_id    891 non-null    int64
1   survived        891 non-null    object
2   pclass          891 non-null    int64
3   name            891 non-null    object
4   sex             891 non-null    object
5   age            714 non-null    float64
6   sib_sp         891 non-null    int64
7   parch          891 non-null    int64
8   ticket         891 non-null    object
9   fare           891 non-null    float64
10  cabin          204 non-null    object
11  embarked       889 non-null    object
dtypes: float64(2), int64(4), object(6)
memory usage: 83.7+ KB
```

```
In [ ]: data['survived'].value_counts()
```

```
Out[ ]: survived
No      549
Yes     342
Name: count, dtype: int64
```

Question 1

```
In [ ]: from sklearn.model_selection import train_test_split
np.random.seed(10)

X = data.drop('survived', axis=1)
Y = data['survived']

x_train, x_test, y_train, y_test = train_test_split(X,Y, train_size=.7, stratify=Y)
```

```
In [ ]: print(f"x_train size: {len(x_train)}, y_train size: {len(y_train)}")
print(f"x_test size: {len(x_test)}, y_test size: {len(y_test)}")

x_train size: 623, y_train size: 623
x_test size: 268, y_test size: 268
```

We can notice below that cabin contains many non-null values, age has over 100, and embarked has 1.

```
In [ ]: x_train.info()

<class 'pandas.core.frame.DataFrame'>
Index: 623 entries, 155 to 253
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   passenger_id    623 non-null   int64  
 1   pclass          623 non-null   int64  
 2   name            623 non-null   object  
 3   sex             623 non-null   object  
 4   age            509 non-null   float64 
 5   sib_sp         623 non-null   int64  
 6   parch          623 non-null   int64  
 7   ticket         623 non-null   object  
 8   fare           623 non-null   float64 
 9   cabin          135 non-null   object  
10   embarked       622 non-null   object  
dtypes: float64(2), int64(4), object(5)
memory usage: 58.4+ KB
```

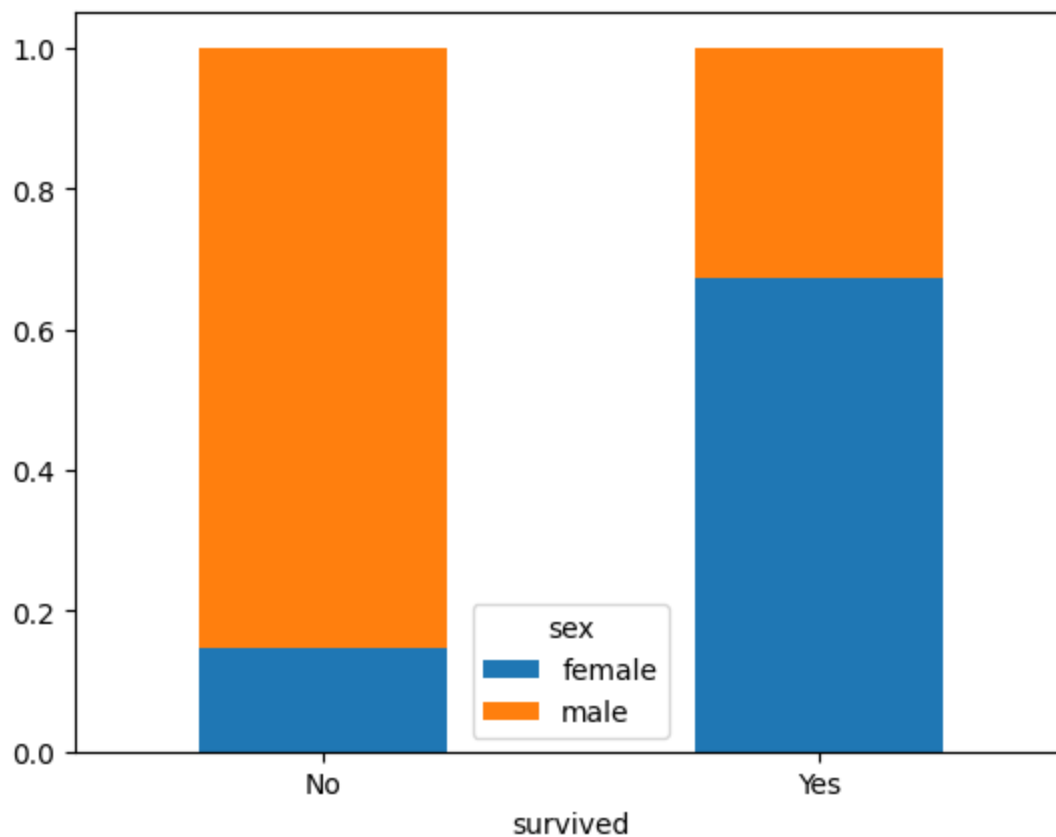
It is a good idea to use stratified sampling for this data since over 60% of passengers don't survive, potentially leading to a class imbalance when the data is split.

Question 2

Yes, it seems sex will be a good predictor of survival outcome, since a much larger proportion of women survived than men.


```
In [ ]: pd.crosstab(index=y_train,
                    columns = x_train['sex'],
                    normalize='index').plot(kind='bar', stacked=True)
```

```
plt.xticks(rotation=0)
plt.show()
```



Yes, it seems that passenger class will be a good predictor of survival outcome, since a large portion of those in classes 1 and 2 survived, while very few in 3rd class survived.

```
In [ ]: pd.crosstab(index=y_train,
                  columns = x_train['pclass'],
                  normalize='index').plot(kind='bar', stacked=True)
plt.xticks(rotation=0)
plt.show()
```



It is useful to use a percent stacked bar chart because of the class imbalance survived and not survived. It also makes the message more clear to a viewer so that they don't have to do any mental calculations.

Question 3

```
In [ ]: # Source: https://seaborn.pydata.org/examples/many\_pairwise\_correlations.html

corr = x_train.get_numeric_data().corr()
mask = np.triu(np.ones_like(corr, dtype=bool))


f, ax = plt.subplots(figsize=(7, 5))

cmap = sns.diverging_palette(230, 20, as_cmap=True)

sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.show()
```

passenger_id -



First, we can see a strong negative correlation between fare and pclass, indicating as expected that a 'lower' class (1 or 2) would be more expensive. Additionally, we also see a moderate, positive correlation between parch and sib_sp, which makes sense given that some parts of a family traveling together makes it likely that others would too. Lastly, there is a noticeable, negative correlation between pclass and age, meaning that older people would more likely to be in a 'lower' class.

Question 4

```
In [ ]: # Remaining features are selected: ticket class, sex, age, number of siblings or spouses
# number of parents or children aboard, and passenger fare
```

```
x_train.drop(['passenger_id', 'name', 'ticket', 'cabin', 'embarked'], axis=1, inplace=True)
x_test.drop(['passenger_id', 'name', 'ticket', 'cabin', 'embarked'], axis=1, inplace=True)
```

```
In [ ]: # Fill in NA values in 'age'
x_train['age'] = x_train['age'].interpolate(method='linear')
x_test['age'] = x_test['age'].interpolate(method='linear')
```

```
In [ ]: # Dummy code categorical variables
x_train = pd.get_dummies(x_train, columns=['pclass', 'sex'], dtype=int, drop_first=True)
x_test = pd.get_dummies(x_test, columns=['pclass', 'sex'], dtype=int, drop_first=True)

x_train.head()
```

```
Out [ ]:
```

| | age | sib_sp | parch | fare | pclass_2 | pclass_3 | sex_male |
|-----|------|--------|-------|---------|----------|----------|----------|
| 155 | 51.0 | 0 | 1 | 61.3792 | 0 | 0 | 1 |
| 14 | 14.0 | 0 | 0 | 7.8542 | 0 | 1 | 0 |
| 421 | 21.0 | 0 | 0 | 7.7333 | 0 | 1 | 1 |
| 432 | 42.0 | 1 | 0 | 26.0000 | 1 | 0 | 0 |
| 484 | 25.0 | 1 | 0 | 91.0792 | 0 | 0 | 1 |

```
In [ ]: # Add interaction terms between sex and fare
x_train['sex_male*fare'] = x_train['sex_male'] * x_train['fare']
x_test['sex_male*fare'] = x_test['sex_male'] * x_test['fare']

# Add interaction terms between age and fare
x_train['age*fare'] = x_train['age'] * x_train['fare']
x_test['age*fare'] = x_test['age'] * x_test['fare']
```

Question 5

```
In [ ]: from sklearn.linear_model import LogisticRegression

# Fit the logistic regression model. Solver chosen to ensure convergence
LR_model = LogisticRegression(solver='newton-cg').fit(x_train, y_train)
```

Question 6

```
In [ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis

# Fit an LDA model
LDA_model = LinearDiscriminantAnalysis().fit(x_train, y_train)
```

Question 7

```
In [ ]: # Fit a QDA Model
QDA_model = QuadraticDiscriminantAnalysis().fit(x_train, y_train)
```

Question 8

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier

# Fit KNN with arbitrary K
KNN_model = KNeighborsClassifier(n_neighbors=10).fit(x_train, y_train)
```

Question 9

Evaluate **Training** Performance

```
In [ ]: predictions = pd.DataFrame([LR_model.predict_proba(x_train)[: ,1],  
                                  LDA_model.predict_proba(x_train)[: ,1],  
                                  QDA_model.predict_proba(x_train)[: ,1],  
                                  KNN_model.predict_proba(x_train)[: ,1]], index=['LR', 'LDA', 'QDA', 'KNN']).T  
predictions.head()
```

```
Out [ ]:
```

| | LR | LDA | QDA | KNN |
|---|----------|----------|----------|-----|
| 0 | 0.410790 | 0.352063 | 0.221037 | 0.8 |
| 1 | 0.700724 | 0.789125 | 0.184691 | 0.1 |
| 2 | 0.161503 | 0.084738 | 0.004210 | 0.1 |
| 3 | 0.693886 | 0.777751 | 0.606448 | 0.5 |
| 4 | 0.375690 | 0.425123 | 0.677802 | 0.6 |

```
In [ ]: from sklearn.metrics import roc_auc_score, roc_curve  
  
for col in predictions.columns:  
    print(f"Area Under ROC curve ({col}): {roc_auc_score(y_train, predictions[col])}")
```

```
Area Under ROC curve (LR): 0.8389938546025104  
Area Under ROC curve (LDA): 0.8402469055090656  
Area Under ROC curve (QDA): 0.8447306485355648  
Area Under ROC curve (KNN): 0.8172561453974896
```

We can see that model performance on training data was very similar for all four models, with KNN performing slightly worse than the other 3.

Question 10

Evaluate **Training** Performance

```
In [ ]: predictions = pd.DataFrame([LR_model.predict_proba(x_test)[: ,1],  
                                  LDA_model.predict_proba(x_test)[: ,1],  
                                  QDA_model.predict_proba(x_test)[: ,1],  
                                  KNN_model.predict_proba(x_test)[: ,1]], index=['LR', 'LDA', 'QDA', 'KNN']).T  
predictions.head()
```

```
Out [ ]:
```

| | LR | LDA | QDA | KNN |
|---|----------|----------|----------|-----|
| 0 | 0.141977 | 0.074279 | 0.003857 | 0.1 |
| 1 | 0.937197 | 0.967739 | 0.999448 | 0.5 |
| 2 | 0.182571 | 0.096587 | 0.004552 | 0.0 |
| 3 | 0.460594 | 0.482569 | 0.262232 | 0.5 |
| 4 | 0.832480 | 0.889982 | 0.778962 | 0.5 |

```
In [ ]: for col in predictions.columns:  
    print(f"Area Under ROC curve ({col}): {roc_auc_score(y_test, predictions[col])}")
```

Area Under ROC curve (LR): 0.8476316563695205
Area Under ROC curve (LDA): 0.8574580759046779
Area Under ROC curve (QDA): 0.7999411591644602
Area Under ROC curve (KNN): 0.7019711679905855


Logistic Regression and LDA had almost identical (testing set) performance based on the Area Under ROC Curve metric. Interestingly, they both performed better on unseen testing data than their training data. This indicates that the assumption of a linear relationship between the log-odds and the features was reasonable, as well as the normal distribution of the features assumption (for LDA). However, we can see that QDA and KNN fit poorly to unseen test data compared to their training data. This indicates that the more flexible models overfit to patterns in the training data that do not exist in the population.

```
In [ ]: from sklearn.metrics import confusion_matrix
```

```
# Make and display confusion matrix  
confusion_matrix(y_test, LDA_model.predict(x_test), normalize='all')
```

```
Out[ ]: array([[0.51492537, 0.10074627],  
              [0.11940299, 0.26492537]])
```

```
In [ ]: y_test_numeric = np.where(y_test == 'Yes', 1, 0)  
  
fpr, tpr, thresholds = roc_curve(y_true=y_test_numeric, y_score=predictions['LDA'])  
  
plt.plot(fpr, tpr)  
plt.plot([0,1], [0,1], ls='--')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve for LDA Model')  
plt.show()
```



Question 11

$$p = \frac{e^z}{1 + e^z} \quad \Rightarrow \quad p + pe^z = e^z$$

$$e^z(1 - p) = p \quad \Rightarrow \quad e^z = \frac{p}{1 - p}$$

$$z = \ln\left(\frac{p}{1 - p}\right)$$

Question 12

$$\begin{aligned}
 a) \quad \text{odds} &= \frac{p}{1-p} = e^z & \Rightarrow \quad \frac{\text{odds}'}{\text{odds}} &= \frac{e^{z'}}{e^z} = e^{\beta_0 + \beta_1(x_1 + z) - (\beta_0 + \beta_1 x_1)} \\
 & & &= e^{2\beta_1} \\
 \text{odds}' &= \frac{p'}{1-p'} = e^{z'} & & \\
 & & \boxed{\text{odds}' = e^{2\beta_1} \cdot \text{odds}} &
 \end{aligned}$$

$$b) \quad p = \frac{1}{1+e^{-z}} \quad z = \beta_0 + \beta_1 x_1$$

$$\text{If } \beta_1 < 0 \quad \text{and} \quad x_1 \rightarrow \infty \quad \Rightarrow \quad z \rightarrow -\infty$$

$$\Rightarrow e^{-z} \rightarrow \infty \quad \Rightarrow \quad p \sim \frac{1}{\infty} \rightarrow 0$$

$$\text{If } \beta_1 < 0 \quad \text{and} \quad x_1 \rightarrow -\infty \quad \Rightarrow \quad z \rightarrow \infty$$

$$\Rightarrow e^{-z} \rightarrow 0 \quad \Rightarrow \quad p \rightarrow 1$$