

Introduction to Statistical Machine Learning

PSTAT 131/231

Winter 2024

INSTRUCTOR:	Dr. Katie Coburn (katie_m_coburn@ucsb.edu)
TEACHING ASSISTANTS:	Miaoqing Yu (miaoqingyu@ucsb.edu) John O'Brien (johnobrien@ucsb.edu) Blythe King (blytheking@ucsb.edu) Qiqi Li (qiqili@ucsb.edu) Hanmo Li (hanmo@ucsb.edu)
LECTURES:	TR 2:00 - 3:15 pm ILP 1101 TR 5:00 - 6:15 pm NH 1006
SECTIONS:	M 10:00 - 10:50 am PHELP 1525 MIAOQING M 11:00 - 11:50 am PHELP 1525 BLYTHE M 12:00 - 12:50 pm PHELP 1526 BLYTHE M 1:00 - 1:50 pm PHELP 1525 QIQI M 2:00 - 2:50 pm PHELP 1526 JOHN M 2:00 - 2:50 pm PHELP 1525 QIQI M 3:00 - 3:50 pm PHELP 1525 JOHN M 4:00 - 4:50 pm PHELP 1513 MIAOQING M 5:00 - 5:50 pm PHELP 1525 HANMO

Course information

Description

STATISTICAL machine learning is used to discover patterns and relationships in large data sets. Topics will include: data exploration, classification and regression trees, random forests, clustering and association rules. Building predictive models focusing on model selection, model comparison and performance evaluation. Emphasis will be on concepts, methods and data analysis; and students are expected to complete a significant class project using real-world data. Prerequisites: PSTAT 120A-B and PSTAT 126 with a minimum grade of C or better. Credit units: 4.

THIS COURSE is taught at two levels, one aimed at undergraduates (131 level) and one at graduate students (231 level). Lectures are given at a level designed to be accessible and relevant to all students in the course. Students taking the course at the 231 level will be assigned **additional homework questions** and must do the reading in **both** the 131 and 231-level textbooks, listed below. The final project assignment is the same for both levels.

Learning Outcomes

Upon completion of this course, students should be able to:

1. EXPLAIN the common statistical learning techniques conceptually and characterize some of them mathematically;
2. DESCRIBE AND DISCUSS the pros and cons of the common statistical learning methods;

3. USE R, the *tidyverse*, and *tidymodels* effectively for exploratory data analysis, model fitting, and visualization.

Format

LECTURES and sections will be provided **in person only**. However, if you are unable to attend due to COVID exposure, symptoms, etc., or for any other university-sanctioned reason, adjustments can be made; please let the instructor know.

Any lecture slides will be posted on Canvas before the corresponding class. The class will progress according to a weekly schedule (provided below). Both lecture periods each week will consist of slides that cover related concepts, sometimes including live coding to demonstrate how to apply the concepts using the R language.

Section each week will be dedicated to completing an R-based lab. Section attendance is not formally required but is **very strongly** recommended, since doing so will enable you to learn and practice your R skills in an environment where you can easily obtain help from the TA(s) and your peers.

Should you need to switch sections, please ask the TA of the section you want to switch to if there is room for you to attend. If the TA allows, you can then begin attending the new section while remaining officially enrolled in your previous section. **Sections will be held beginning the week of January 15th.**

Required Materials

THE BOOKS and software used for this class are as follows. Note that all are freely available online and do not require purchase.

- **Required for PSTAT 131 students:** *An Introduction to Statistical Learning with Applications in R* (or ISLR), by G. James, D. Witten, T. Hastie, R. Tibshirani. Available: <https://www.statlearning.com/>
- **Required for PSTAT 231 students:** *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (or ESL), by T. Hastie, R. Tibshirani, J. Friedman. Available: <https://hastie.su.domains/ElemStatLearn/>
- **Required for all:**
 - The R statistical environment version 4.3.2, available here: <https://www.r-project.org/>
 - The RStudio IDE (integrated development environment) Desktop version. Available: <https://www.rstudio.com/>

Supplemental Materials

THE FOLLOWING are materials that are **not required** for this course, but are **recommended**. Don't feel expected to read all of these; I encourage you to use them as reference manuals and look subjects up as necessary.

- *R for Data Science*, by H. Wickham, G. Grolemund. Available: <https://r4ds.had.co.nz/index.html>
- *The StatQuest Illustrated Guide to Machine Learning*, by Josh Starmer. Available on Amazon: \$27 for paperback. Based on videos available free at <https://www.youtube.com/@statquest>
- *Supervised Machine Learning for Text Analysis in R*, by E. Hvitfeldt and J. Silge. Available: <https://smltar.com/>
- *Feature Engineering and Selection*, by M. Kuhn and K. Johnson. Available: <http://www.feats.engineering/>
- *R Markdown Cookbook*, by Y. Xie, C. Dervieux and E. Riederer. Available: <https://bookdown.org/yihui/rmarkdown-cookbook/>
- *Tidy Modeling with R*, by M. Kuhn and J. Silge. Available: <https://www.tmwr.org/>

Assessments

YOUR ATTAINMENT of course learning outcomes will be measured by the following assessments, with the relative weighting indicated in parentheses. All assignments within each category are given equal weight.

Collaboration and open science are fundamental aspects of data science. You can (and in fact are encouraged) to use the Internet, StackOverflow, etc. as a resource when coding. However, you **should not** directly copy code or answers. Your code should be well documented; include comments describing each line. If you use a source for help coding, you **must** cite it. You **are permitted** to collaborate/work in groups on the homework if you choose, but everyone should work on their own final project – it is not a group assignment.

You will **not** be graded on your completion of lab work or your attendance of labs, but doing the labs will be greatly beneficial for your homework and final project grades. You are strongly encouraged to attend lab sections and follow along with the materials for each lab, which will be posted each week.

- **Homework** (50%). Homework is assigned approximately every two weeks (this may vary) and will always be due by Sundays at 11:59 PM PDT. The graduate students will need to complete additional questions, but all assignments are ultimately rescaled to be out of 50 points. There are a total of **five** homework assignments and none will be dropped.
- **Final Project** (50%). The final project is a significant portion of the class and allows students to demonstrate their understanding of the material by working with a data set of their own choosing. Each student should begin thinking about their project and looking for a data set immediately; a page on Canvas provides a list of possible resources. You will receive a lot of guidance about the project over the course of the quarter.

The project will require a large amount of work and cannot be satisfactorily completed in only a day (or even a week). You are **strongly advised** to work on it in stages throughout the quarter. See Canvas for more detailed instructions, a rubric, and several examples of exemplary past projects.

Course Schedule

THE WEEKLY schedule is indicated below. The topics and reading are subject to change based on the progress of the class.

	Topic	Assignment	Project Stage
Week 1	Intro to ML	Data Memo	Pick topic
Week 2	Linear models	Homework 1	Find data
Week 3	Logistic models		Tidy data
Week 4	Discriminant analysis	Homework 2	EDA
Week 5	Resampling & feature engineering		
Week 6	Regularization	Homework 3	Run models
Week 7	Tree-based methods	Homework 4	
Week 8	Clustering		Write-up
Week 9	SVMs	Homework 5	Edits
Week 10	Neural networks	Final Project	Final draft

Time Commitment

THE COURSE is 4 credit units; each credit unit corresponds to an approximate time commitment of 3 hours. You should expect to allocate 12 hours per week to the course. If you find yourself spending considerably more time on the course on a regular basis, please let the instructor or TAs know so that we can help you balance the workload.

A suggested allocation of this time is as follows:

- Reading and class time: 3 hours (25%)
- Homework: 4.5 hours (37.5%)
- Sections: 1 hour (8.3%)
- Final project: 4.5 hours (37.5%)

Course Policies

Communication

THERE ARE FOUR means of communication with other students or the instructional team: during/after class, office hours, email, and individual appointments. **Please use them in that order.** Note that I can be slow at responding to email, which is why it is low on the list; to compensate, however, I offer several office hours weekly and am always willing to discuss problems or answer questions in person/during class.

1. **During/after class.** The easiest guaranteed way to contact me is to come up after class and say hello, or to raise your hand and ask a question during class.
2. **Office hours.** Office hours are offered at a minimum of twice weekly. I hold one session of 4 hours a week, and the other(s) are held by various combinations of TAs or ULAs. These are opportunities to interact informally, ask questions, and discuss course material or assignments.
3. **Email.** Please use email with discernment for simple communication. A response is guaranteed within 48 weekday hours (so if you email on Friday afternoon, you may not receive a reply until Tuesday afternoon). If your message is time-sensitive, please indicate so in the subject, and I will do my best to respond promptly.
4. **Appointment.** You can schedule individual 20-minute appointments with me as needed. These appointments may be either on Zoom or in person. This mode of communication is best suited to more complex or nuanced communication regarding personal matters. If you schedule an appointment, you will be prompted to indicate what you wish to discuss.

Extra Credit

IF THE CLASS reaches a 90% submission rate of ESCI surveys at the end of the quarter, the **entire class** will receive 2 free points on the final project.

Grades

YOUR OVERALL GRADE in the course will be calculated as the weighted average of the proportions of total possible points in each assessment category according to the weightings indicated in the Assessments section and reported as a percentage rounded to two decimal places. Letter grades will be assigned according to the rubric below.

A+	100%
A	94% – 99.99%
A-	90% – 93.99%
B+	87% – 89.99%
B	84% – 86.99%
B-	80% – 83.99%
C+	77% – 79.99%
C	74% – 76.99%
C-	70% – 73.99%
D+	67% – 69.99%
D	64% – 66.99%
D-	60% – 63.99%
F	0% – 59.99%

YOU CAN keep track of your marks on individual assessments, your marks in each assessment category, and your overall grade in the Canvas gradebook. Please notify the instructor or TAs of any errors in grade *entry*; please do not attempt to negotiate the grades themselves. If at the end of the course you believe your grade was unfairly assigned, you are entitled to contest it according to the procedure outlined here in the UCSB General Catalog.

Deadlines

YOU RECEIVE two free late homework submissions without penalty. This policy applies only to homeworks. When you wish to use one of these late submissions, simply submit within one week of the original deadline, and **add a note** at the beginning of the assignment explaining that it is a free late submission.

NON-EXEMPTED homeworks submitted within one week of the deadline will be awarded 50% credit. **No credit** will be awarded for homework turned in more than one week late; please plan ahead and submit your work on time.

THE FINAL PROJECT deadline is firm and no late submissions will be accepted.

Extensions

EXTENSIONS may be granted based on individual circumstances at the instructor's discretion.

Conduct

PLEASE BE MINDFUL of maintaining respectful and kind communication. You are expected to uphold the UCSB student code of conduct in your behavior when in class, in section, or interacting with other students or the instructional team. You can find the student code of conduct on the Office of Student Conduct website from this page. If you are uncomfortable with the conduct of another individual for any reason, please notify the instructor or TAs.

Academic Integrity

PLEASE MAINTAIN INTEGRITY in learning. Your work in the course must be your own. Any form of plagiarism, cheating, misrepresentation of individual effort on assignments and assessments, falsification of information or documents, or misuse of course materials compromises your own learning experience, that of your peers, and undermines the integrity of the UCSB community. Any evidence of dishonest conduct will be discussed with the student(s) involved and reported to the Office of Student Conduct. Depending on the nature of the evidence and the violation, penalty in the course may range from loss of credit to automatic failure. For a definition and examples of dishonesty, a discussion of what constitutes an appropriate response from faculty, and an explanation of the reporting and investigation process, see the OSC page on academic integrity.

Accommodations

REASONABLE ACCOMMODATIONS will be made for any student with a qualifying disability. Such requests should be made through the Disabled Students Program (DSP). More information, instructions on how to access accommodations, and information on related resources can be found on DSP website. Remote learning may present unique accommodation needs requiring additional flexibility; students receiving accommodation via DSP are invited to discuss this with the instructor if desired.

Student Evaluations

TOWARD THE END of the term, you will be given an opportunity to provide feedback about the course. Your suggestions and assessments are essential to improving the course, so please take the time to fill out the evaluations thoughtfully.

Student Resources

ANY STUDENTS in need are encouraged to make use of the following resources.

- Financial Crisis Response Team
<https://food.ucsb.edu/about/committees/financial-crisis-response-team>
- Food Security and Basic Needs (Food, Housing, Technology) Advising Center
<https://food.ucsb.edu/resources/basic-needs-advocates>
- Undocumented Student Services
<http://www.sa.ucsb.edu/dreamscholars/home>
- Campus Advocacy, Resources, and Education (CARE)
<https://care.ucsb.edu/>
24/7 Confidential Phone: (805) 893-4613
- The Trevor Project
<https://www.thetrevorproject.org/>
- Counseling and Psychological Services (CAPS)
<https://caps.sa.ucsb.edu/>
24/7 Counselors: (805) 893-4411, press 2